





Developing a Linux driver in Rust for ARM




Who are we ?





RFID-RC522 pinout





Why Rust?

- Safety
- Speed
- Concurrency
- Expressivity
- No garbage collection and no manual memory management! (lifetime analysis)
- Zero-cost abstractions
- Zero setup cross compilation
- Strong ecosystem and tons of libraries at your fingertips!
- **Not for rewriting the kernel!**



First method: Calling Rust from C



```
// mfr522.c
void print(char *msg) {
    printk(KERN_INFO "%s", msg);
}

static int __init mfr522_init(void) {
    hello_rust();
    return 0;
}

static void __exit mfr522_exit(void) {
    pr_info("MFRC522-Rust exit\n");
}

module_init(mfr522_init);
module_exit(mfr522_exit);
```

```
// mfr522.rs
#![no_std]

extern "C" {
    fn print(msg: *const u8) -> ();
}

pub extern "C" fn hello_rust() -> i32 {
    unsafe {
        print("Hello from Rust\n\0".as_bytes().as_ptr())
    };

    return 0;
}
```



C glue for Rust code

Pros

- It works (kinda...)
- 2 ways to go about it:
 - C -> Rust can't "talk" with the kernel (compute only)
 - C -> Rust -> C can "talk" with the kernel

Cons

- Needs a C to Rust wrapper for each kernel function



```
// .cargo/config.toml
[build]
rustflags = [
    "-C", "relocation-model=static", "-C", "code-model=kernel", "-Z", "plt=y"
]

// Works... But "code-model=kernel" is only available for x86, so we removed it...
```




Issues with the method

- Weird requirements from the kernel's *modpost* tool
 - Issues with unwinding functions defined in *libgcc*



```
// Cargo.toml  
[profile.dev]  
panic = "abort"  
  
[profile.release]  
panic = "abort"
```



```
ERROR: modpost: "__aeabi_unwind_cpp_pr1" [mfrc522.ko] undefined!
```

```
ERROR: modpost: "__aeabi_unwind_cpp_pr0" [mfrc522.ko] undefined!
```

```
// `cpp`? In my Rust code?
```

```
// Stack unwinding? But we abort on panics!
```



Issues with the method

- Common issue when cross-compiling for ARM?
- Since we don't need unwinding code, we can write those functions ourselves, right?
 - More glue!



```
// If we arrive here, we're screwed anyway  
#[no_mangle]  
unsafe extern "C" fn __aeabi_unwind_cpp_pr1() -> ! { panic!() }  
#[no_mangle]  
unsafe extern "C" fn __aeabi_unwind_cpp_pr0() -> ! { panic!() }
```



Issues with the method

- *modpost* does not complain...
- ...but *insmod* does, missing relocation
- Only present when using Rust features: *pattern matching*, *traits*, ...
 - Which defeats the purpose of using Rust...



Second method: Using rust-for-linux



What is Rust-for-Linux?



What is Rust for Linux?

- <https://github.com/Rust-for-Linux/linux>
- rust-for-linux@vger.kernel.org
- Goal: add support for the Rust language to the Linux kernel
- The code will be submitted to review to LKML
- First patch to linux-next: <https://lkml.org/lkml/2021/4/14/1023>



Compile rust-for-linux/linux

```
$ rustup default nightly-2021-02-20
...
$ rustup component add rust-src
...
$ cargo install --locked --version 0.56.0 bindgen
...
$ # Enable Rust support: General setup -> Rust support
$ make ARCH=arm64 CROSS_COMPILE=aarch64-linux-gnu- LLVM=1 LLVM_IAS=1
...
```



Out of tree module in Rust for Linux

- Compiles correctly for arm64 defconfig...
- ...but not for our Raspberry Pi 3b+ config



```
error[E0463]: can't find crate for `core`
```

```
|
```

```
= note: the `target-10773475644222393364` target may not be installed
```

```
error: aborting due to previous error
```

For more information about this error, try `rustc --explain E0463`.

```
// target-10773475644222393364??
```

```
// Seems to be a cross-compilation issue
```



Rust for Linux, native
compilation

- We could use a VM, but that's heavy and not practical
- What about using a RasPi?
- Compilation time will increase
- What about chrooting? We do not need another running kernel



Chrooting in an aarch64 rootfs

- Impossible ?
- Use *qemu-user-static*
- Download a rootfs
- *chroot* into it
- Et voilà!

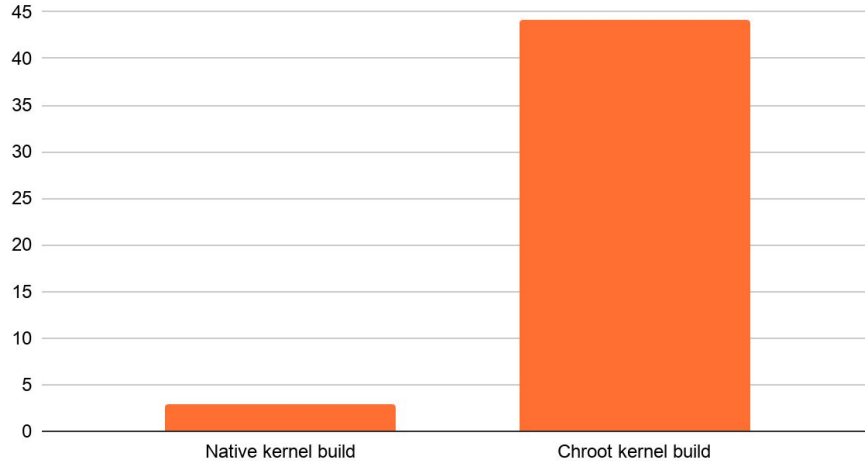


Compiling in the chroot

- Sloooooooooow
- "Incremental" build time on chroot ~= full native build time
- Workflow less than pleasant.
- But out of tree module compilation works so...
- But still faster than on a Raspberry Pi
- After some time we switched to intree compilation



Native VS chroot compilation time (minutes)



```
1 [|||||] [98.2%] 21 [|||||] [100.0%] 41 [|||||] [94.5%] 61 [|||||] [88.3%]
2 [|||||] [100.0%] 22 [|||||] [94.5%] 42 [|||||] [95.1%] 62 [|||||] [91.4%]
3 [|||||] [100.0%] 23 [|||||] [98.8%] 43 [|||||] [98.8%] 63 [|||||] [96.9%]
4 [|||||] [99.4%] 24 [|||||] [100.0%] 44 [|||||] [100.0%] 64 [|||||] [98.2%]
5 [|||||] [100.0%] 25 [|||||] [100.0%] 45 [|||||] [100.0%] 65 [|||||] [100.0%]
6 [|||||] [98.8%] 26 [|||||] [100.0%] 46 [|||||] [100.0%] 66 [|||||] [98.2%]
7 [|||||] [97.5%] 27 [|||||] [100.0%] 47 [|||||] [100.0%] 67 [|||||] [100.0%]
8 [|||||] [100.0%] 28 [|||||] [100.0%] 48 [|||||] [98.2%] 68 [|||||] [98.8%]
9 [|||||] [95.7%] 29 [|||||] [98.8%] 49 [|||||] [100.0%] 69 [|||||] [100.0%]
10 [|||||] [98.8%] 30 [|||||] [100.0%] 50 [|||||] [94.4%] 70 [|||||] [99.4%]
11 [|||||] [94.4%] 31 [|||||] [100.0%] 51 [|||||] [100.0%] 71 [|||||] [100.0%]
12 [|||||] [98.8%] 32 [|||||] [100.0%] 52 [|||||] [100.0%] 72 [|||||] [99.4%]
13 [|||||] [100.0%] 33 [|||||] [100.0%] 53 [|||||] [96.3%] 73 [|||||] [100.0%]
14 [|||||] [100.0%] 34 [|||||] [100.0%] 54 [|||||] [93.9%] 74 [|||||] [100.0%]
15 [|||||] [99.4%] 35 [|||||] [100.0%] 55 [|||||] [100.0%] 75 [|||||] [97.6%]
16 [|||||] [96.9%] 36 [|||||] [100.0%] 56 [|||||] [100.0%] 76 [|||||] [100.0%]
17 [|||||] [100.0%] 37 [|||||] [96.3%] 57 [|||||] [100.0%] 77 [|||||] [98.2%]
18 [|||||] [100.0%] 38 [|||||] [98.2%] 58 [|||||] [99.4%] 78 [|||||] [100.0%]
19 [|||||] [100.0%] 39 [|||||] [100.0%] 59 [|||||] [96.9%] 79 [|||||] [100.0%]
20 [|||||] [100.0%] 40 [|||||] [100.0%] 60 [|||||] [100.0%] 80 [|||||] [97.0%]
Mem[|||||] [3.66G/98.2G]
Swp[|||||] [6.01M/8.00G]
Tasks: 243, 196 thr; 78 running
Load average: 8.24 2.61 1.74
Uptime: 28 days, 03:31:34
```




Writing the driver



What do we need ?

- MODULE_{LICENSE, AUTHOR, DESCRIPTION} ✓
- module_{init, exit} ✓
- printk ✓
- misc devices API ✓
- spi API ✗



Rust API generation with bindgen

```
// rust/kernel/bindings_helper.h
```

```
/* SPDX-License-Identifier: GPL-2.0 */  
#include <linux/cdev.h>  
#include <linux/fs.h>  
#include <linux/module.h>  
#include <linux/random.h>  
#include <linux/slab.h>  
...  
#include <uapi/linux/android/binder.h>
```

```
// rust/bindings_generated.rs
```

```
pub const PLATFORM_DEVID_NONE: i32 = -1;  
...  
pub struct sysinfo {  
    pub uptime: __kernel_long_t,  
    ...  
}  
...  
extern "C" {  
    pub fn misc_register(misc: *mut  
miscdevice) -> c_types::c_int;  
}
```



Rust SPI API

- Add `#include <kernel/spi.h>` to `rust/kernel/bindings_helper.h`
- We now have generated Rust bindings for `spi_register`, `spi_write_then_read` and so on
- But our work is not done!
 - Simply using C functions from Rust negates the advantages of Rust
 - Rust-for-Linux implements high-level abstractions for interacting with the kernel
 - So we need to implement our abstraction too!



```
static struct spi_driver mfrc522_spi_driver = {
    .driver = {
        .name = "mfrc522", .owner = THIS_MODULE,
    },
    .probe = mfrc522_spi_probe,
};

static int mfrc522_spi_probe(struct spi_device *client) {}

static int __init mfrc522_init(void) {
    int ret = spi_register_driver(&mfrc522_spi_driver);
    if (ret) {
        pr_err("[MFRC522] SPI Register failed\r\n");
        return ret;
    }
    return 0;
}
```



```
struct Mfrc522SpiMethods;

impl SpiMethods for Mfrc522SpiMethods {
    declare_spi_methods!(probe);

    fn probe(mut spi_device: SpiDevice) -> Result {}
}

impl KernelModule for Mfrc522Driver {
    fn init() -> Result<Self> {
        let spi = spi::DriverRegistration::new_pinned::<Mfrc522SpiMethods>(
            &THIS_MODULE,
            cstr!("mfrc522"),
        )?;
    }
}
```



How do we achieve this?

- Allow the user to register an SPI Driver (and unregister it safely)
- Use high level abstractions for the definitions of SPI Methods
- Add wrappers around usual SPI functions
 - *spi_write_then_read*
 - *spi_write*
 - *spi_read*
 - ...
- The user should only interact with safe, idiomatic Rust code



```
// Generated Rust representation of the spi_driver struct
```

```
#[repr(C)]  
#[derive(Debug, Copy, Clone)]  
pub struct spi_driver {  
    ...  
    pub probe: Option<unsafe extern "C" fn(spi: *mut spi_device) -> c_types::c_int>,  
    pub remove: Option<unsafe extern "C" fn(spi: *mut spi_device) -> c_types::c_int>,  
    pub shutdown: Option<unsafe extern "C" fn(spi: *mut spi_device)>,  
    ...  
}
```




```
// The `probe` method is an option: Equivalent to NULL or a function
// It should take a pointer as argument, and return an integer
pub struct spi_driver {
    pub probe: Option<unsafe extern "C" fn(spi: *mut spi_device) -> c_types::c_int>,
}

// Equivalent to C's
struct spi_driver {
    int (*probe)(struct spi_device *spi);
}
```



Why external functions are unsafe

- Since the compiler cannot guarantee the safety of external C functions, calling them is inherently unsafe
- We need to make sure to document unsafe code as well as check it carefully, to avoid undefined behavior
 - Which we encountered... But more on that later



```
// We could have the user write the following
unsafe extern "C" fn my_probe(spi: *mut spi_device) -> c_types::c_int {
    pr_info!("Probing successful!\n");

    0
}
```



High level SPI methods

- But we can't let the user write unsafe functions!
- What's the point of using Rust if it's just to interact with C?
- It makes the user interact with raw pointers and C types
 - This is quite annoying in Rust!
 - Rust is written using references instead of pointers
 - And using *i32* instead of *c_types::c_int*
 - *c_types* are FFI safe and correspond to the actual C types and API, but they're still annoying to use



// Ideally, this is what `probe` should look like

```
fn my_probe(mut spi: SpiDevice) -> Result {  
    pr_info!("Probing successful!\n");  
  
    Ok(())  
}
```

// The function is safe

// We use high level types such as `SpiDevice`

// It returns using a high level abstraction, `Results`, instead of integers



```
// Here's what ours look like (cropped a little)
fn probe(mut spi_device: SpiDevice) -> Result {
    pr_info!("[MFRC522-RS] SPI Registered\n");

    let version = match Mfrc522Spi::get_version(&mut spi_device) {
        Ok(v) => v,
        Err(_) => return Err(kernel::Error::from_kernel_errno(-1)),
    };

    pr_info!("[MFRC522-RS] MFRC522 {:?} detected\n", version);

    Ok(())
}
```



Converting safe, high-level functions into low level ones

- Since we interact with C, we cannot pass our high-level function as a member to our `spi_driver`. We *need* to give it an unsafe function using C types and pointers.
- We can convert a Rust function into a C function using Rust macros! (Don't do this)



```
macro_rules! spi_method {
    (fn $method_name:ident (mut $device_name:ident : SpiDevice) -> Result $block:block) => {
        unsafe extern "C" fn $method_name(dev: *mut kernel::bindings::spi_device) ->
kernel::c_types::c_int {
            use kernel::spi::SpiDevice;

            fn inner(mut $device_name: SpiDevice) -> Result $block

            match inner(SpiDevice::from_ptr(dev)) {
                Ok(_) => 0,
                Err(e) => e.to_kernel_errno(),
            }
        }
    };
}
```




```
// This is the invocation, what you use to call the macro
(fn $method_name:ident (mut $device_name:ident : SpiDevice) -> Result $block:block)
//  ~~~~~ ~~~~~ ~~~~~
// These are our macro parameters
// The rest are just matching keywords for invoking the macro!
```



```
// We could also use this match arm, for example  
(penguins $method_name:ident love $device_name:ident crabs $block:block)
```



```
// This is the generated code. $Arguments are expanded from the invocation
// This expansion is done at the AST level, not as tokens in the source
unsafe extern "C" fn $method_name(dev: *mut kernel::bindings::spi_device) ->
kernel::c_types::c_int {
    use kernel::spi::SpiDevice;

    fn inner(mut $device_name: SpiDevice) -> Result $block

    match inner(SpiDevice::from_ptr(dev)) {
        Ok(_) => 0,
        Err(e) => e.to_kernel_errno(),
    }
}
```



```
macro_rules! spi_method {
    (fn $method_name:ident (mut $device_name:ident : SpiDevice) -> Result $block:block) => {
        unsafe extern "C" fn $method_name(dev: *mut kernel::bindings::spi_device) ->
        kernel::c_types::c_int {
            use kernel::spi::SpiDevice;

            fn inner(mut $device_name: SpiDevice) -> Result $block

            match inner(SpiDevice::from_ptr(dev)) {
                Ok(_) => 0,
                Err(e) => e.to_kernel_errno(),
            }
        }
    };
}
```



// Let's use our macro for our previously defined probe

```
fn my_probe(mut spi: SpiDevice) -> Result {  
    pr_info!("Probing successful!\n");  
  
    Ok(())  
}
```



```
// Let's use our macro for our previously defined probe
```

```
spi_method! {  
    fn my_probe(mut spi: SpiDevice) -> Result {  
        pr_info!("Probing successful!\n");  
  
        Ok(())  
    }  
}
```

```
// Done!
```



```
// Or the alternative...
spi_method! {
    penguins my_probe love spi crabs {
        pr_info!("Probing successful!\n");

        Ok(())
    }
}
```



// Let's look at the generated code

```
unsafe extern "C" fn my_probe(dev: *mut kernel::bindings::spi_device) -> kernel::c_types::c_int {  
    use kernel::spi::SpiDevice;  
  
    fn inner(mut spi: SpiDevice) -> Result {  
        pr_info!("Probing successful!\n");  
  
        Ok(())  
    }  
  
    match inner(SpiDevice::from_ptr(dev)) {  
        Ok(_) => 0,  
        Err(e) => e.to_kernel_errno(),  
    }  
}
```




Probe function, macro edition

Pros

- Easy to implement
- Easy to use

Cons

- Unmaintainable
- *Awful* compiler errors on misuse



What's the alternative

- We can use Rust's object paradigm to create a *trait*
 - This defines an interface which a type can implement



// This is similar to Java interfaces for example. This one has default implementations

```
pub trait SpiMethods {  
    fn probe(mut _spi_dev: SpiDevice) -> Result {  
        Ok(())  
    }  
  
    fn remove(mut _spi_dev: SpiDevice) -> Result {  
        Ok(())  
    }  
  
    fn shutdown(mut _spi_dev: SpiDevice) {}  
}
```



What's the alternative

- Using genericity, we can implement functions with the correct valid signature
- Which take any high level Rust type



```
// This function has the right signature for a probe function, but is generic!  
// Will generate a function specific to this type  
// We can have trait bounds (T: SpiMethods) to make sure that the types we use it on implement the  
// required methods  
unsafe extern "C" fn probe_wrapper<T: SpiMethods>(s: *mut bindings::spi_device) -> c_types::c_int {  
    // We call the probe function from the type given as generic parameter  
    match T::probe(SpiDevice::from_ptr(s)) {  
        Ok(_) => 0,  
        Err(e) => e.to_kernel_errno(),  
    }  
}
```



```
// The register method is generic too, with a trait bound too
pub fn register<T: SpiMethods>(self: Pin<&mut Self>) -> Result {
    // Since our wrappers have the right signature, we can just give them to our low level
    // spi_driver!
    this.spi_driver.probe = Some(DriverRegistration::probe_wrapper::<T>());
    this.spi_driver.remove = Some(DriverRegistration::remove_wrapper::<T>());
    this.spi_driver.shutdown = Some(DriverRegistration::shutdown_wrapper::<T>());
}
```

What's the alternative

- However, we can't use the kernel's default implementations, since we always pass functions and never NULL
- We need a way for the user to specify when they want to implement a function, and when to use the Kernel's default
- We could use reflection: If the type T implements the function *probe*, use it, otherwise use NULL... This is really slow!
- We can use a virtual/truth table by having the user declare what functions they will implement. This is what's being done for the *FileOperations* abstraction



```
pub struct ToUse {  
    pub probe: bool,  
    pub remove: bool,  
    pub shutdown: bool,  
}
```

```
// We can define constants inside of a trait for implementers to use
```

```
pub trait SpiMethods {  
    const TO_USE: ToUse;  
  
    ...  
}
```

```
// Since it's not really pretty nor interesting, we can wrap the declaration of this constant  
// in a macro
```




```
impl SpiMethods for Mfrc522SpiMethods {  
    declare_spi_methods!(probe);
```

```
    ...
```

```
    // Remember to implement probe properly, or the trait's default implementation (not the  
// kernel's) will be used. This implementation basically does nothing.  
// This macro basically declares the `TO_USE` constant with all fields set to false,  
// except `TO_USE.probe` which is set to true
```

```
}
```



```
pub fn register<T: SpiMethods>(self: Pin<&mut Self>) -> Result {
    this.spi_driver.probe = if T::TO_USE.probe {
        Some(DriverRegistration::probe_wrapper::<T>)
    } else {
        None
    };
    this.spi_driver.remove = if T::TO_USE.remove {
        Some(DriverRegistration::remove_wrapper::<T>)
    } else {
        None
    };
    // Ugly code, but this is just from the low-level implementation. The user only deals with
    // idiomatic Rust.
}
```



```
pub fn register<T: SpiMethods>(self: Pin<&mut Self>) -> Result {  
    this.spi_driver.probe = T::TO_USE.probe.then_some(DriverRegistration::probe_wrapper::<T>);  
  
    // Rust has us covered!  
}
```



Probe function, trait edition

Pros

- Better type safety

Cons

- Harder to understand and implement



```
pub fn register(self: &mut Self) -> Result {
    let mut spi_driver =
bindings::spi_driver::default();
    ...
    // Here spi_driver is copied...
    self.spi_driver = Some(spi_driver);
    let res = unsafe {
        // Inside an unsafe block, rustc can't guarantee
        // that spi_driver will not be used outside the
        // scope of the current function
bindings::__spi_register_driver(this.this_module.0,
&mut spi_driver)
    };
    ...
}
```

```
[ 36.266704] mfr522: [MFR522-RS] Init
[ 36.267087] mfr522: [MFR522-RS] SPI Registered
[ 36.272868] Unable to handle kernel paging request at virtual address ffff80001087b7d0
[ 36.282778] Mem abort info:
[ 36.286541]   ESR = 0x96000007
[ 36.290281]   EC = 0x25: DABT (current EL), IL = 32 bits
[ 36.296322]   SET = 0, FnV = 0
[ 36.300027]   EA = 0, S1PTW = 0
[ 36.303836] Data abort info:
[ 36.307375]   ISV = 0, ISS = 0x00000007
[ 36.311887]   CM = 0, WnR = 0
[ 36.315514] swapper pgtable: 4k pages, 48-bit VAs, pgdp=0000000013f1000
[ 36.322967] [ffff80001087b7d0] pgd=0000000001810003, p4d=0000000001810003, pud=0000000001811003, pmd=00000000
[ 36.337034] Internal error: Oops: 96000007 [#1] PREEMPT SMP
[ 36.343378] Modules linked in: mfr522 hci_uart btqca btbcm bluetooth ecdh_generic ecc 8021q mrp garp stp llc
[ 36.378161] CPU: 3 PID: 539 Comm: systemd-udevd Tainted: G               W               5.12.0-rc4+ #25
[ 36.388189] Hardware name: Raspberry Pi 3 Model B Plus Rev 1.3 (DT)
[ 36.395343] pstate: 80000005 (Nzcv daif -PAN -UAO -TCO BTYP=--)
[ 36.402250] pc : dev_uevent+0x138/0x1e8
[ 36.406966] lr : uevent_show+0x90/0x118
[ 36.411677] sp : ffff80001079bba0
[ 36.415849] x29: ffff80001079bbc0 x28: ffff670b41b9d400
[ 36.422066] x27: 0000000000000000 x26: 0000000000000001
[ 36.428281] x25: 0000000000000000 x24: ffff670b49b42700
[ 36.434496] x23: ffff670b41b9d400 x22: ffff670b41b9d400
[ 36.440711] x21: ffff670b42475800 x20: ffff670b47f59000
[ 36.446900] x19: ffff670b42475800 x18: 0000000000000000
[ 36.453060] x17: 0000000000000000 x16: 0000000000000000
[ 36.459216] x15: ffff670b4342f0d0 x14: 0000000000000000
[ 36.465374] x13: 0000000000000001 x12: 0000000000000000
[ 36.471522] x11: 0000000000000001 x10: ffff670b47f59000
[ 36.477660] x9 : 0de4e88efe7dd600 x8 : ffff80001087b7d0
[ 36.483806] x7 : 0000000000000000 x6 : 000000000000003f
[ 36.489957] x5 : 0000000000000040 x4 : ffff80001079bb00
[ 36.496107] x3 : 0000000000000001 x2 : ffff670b47f59000
[ 36.502254] x1 : ffff670b42475800 x0 : ffff670b41b9d400
[ 36.508389] Call trace:
[ 36.511582]   dev_uevent+0x138/0x1e8
[ 36.515824]   uevent_show+0x90/0x118
[ 36.520048]   dev_attr_show+0x20/0x58
[ 36.524339]   sysfs_kf_seq_show+0xa0/0x110
[ 36.529046]   kernfs_seq_show+0x2c/0x9c
[ 36.533465]   seq_read_iter+0x11c/0x3b4
[ 36.537858]   kernfs_fop_read_iter+0x68/0x188
[ 36.542758]   vfs_read+0x290/0x2bc
[ 36.546666]   ksys_read+0x74/0xe0
[ 36.550464]   __arm64_sys_read+0x1c/0x28
[ 36.554867]   el0_svc_common+0x90/0x110
[ 36.559170]   do_el0_svc+0x24/0x80
[ 36.563029]   el0_svc+0x28/0x88
[ 36.566622]   el0_sync_handler+0x84/0xe4
[ 36.571006]   el0_sync+0x154/0x180
[ 36.574856] Code: aa1403e0 97f8c37e f9403668 b40000c8 (f9400102)
[ 36.581549] ---[ end trace 556d98a75f645ca9 ]---
```



Rust for Linux

Pros

- It works!

Cons

- Extremely recent
- Subject to changes
- A lot of work still needs to be done



MFRC522 driver: C vs Rust



```
enum Mfrc522Command {  
    Idle = 0,  
    Mem  
};  
  
void test(enum Mfrc522Command e) {  
    printf("%i\n", e);  
}  
  
int main(void) {  
    test(Idle);  
    test(42);  
}
```

// Compiles fine

```
enum Mfrc522Command {  
    Idle = 0b0000,  
    Mem = 0b0001,  
}  
  
fn test(e: Mfrc522Command) {  
    println!("{:?}", e);  
}  
  
fn main() {  
    test(Mfrc522Command::Idle);  
    test(42);  
}
```

```
error[E0308]: mismatched types  
--> src/main.rs:12:10  
    |  
12  |         test(42);  
    |         ^^ expected enum `Mfrc522Command`,  
found integer
```




```
struct driver_command {
    const char *input; u8 cmd;
};

struct driver_command commands[MFRC522_CMD_AMOUNT] = {
    { .input = "mem_write", .cmd = MFRC522_CMD_MEM_WRITE },
    ...
    { .input = "version", .cmd = MFRC522_CMD_GET_VERSION },
    { .input = "debug", .cmd = MFRC522_CMD_DEBUG },
};

struct driver_command *find_cmd_from_token(const char *token)
{
    size_t i;

    for (i = 0; i < MFRC522_CMD_AMOUNT; i++)
        if (!strcmp(token, commands[i].input))
            return &commands[i];

    return NULL;
}
```

```
pub enum Cmd {
    MemWrite,
    MemRead,
    GetVersion,
    GenRand,
}

impl Cmd {
    pub fn from_str(cmd: &str) -> Option<Cmd> {
        match cmd {
            "mem_read" => Some(Cmd::MemRead),
            "mem_write" => Some(Cmd::MemWrite),
            "get_version" => Some(Cmd::GetVersion),
            "gen_rand_id" => Some(Cmd::GenRand),
            _ => None,
        }
    }
}
```



```
static inline int
spi_read(struct spi_device *spi,
         void *buf, size_t len)
{
    struct spi_transfer t = {
        .rx_buf = buf,
        .len = len,
    };

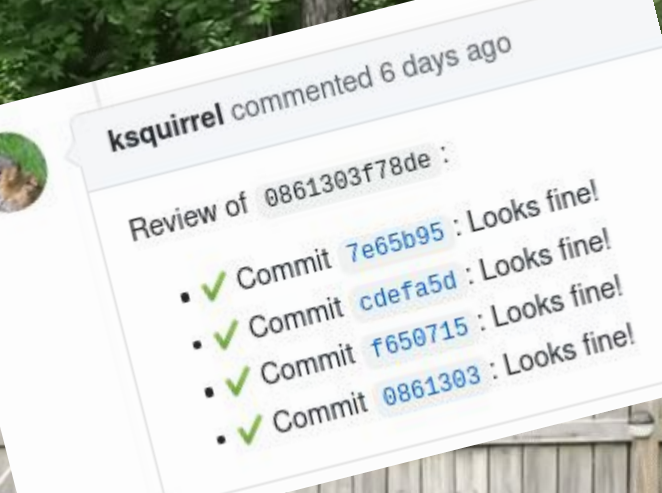
    return spi_sync_transfer(spi, &t, 1);
}
```

```
#[inline]
fn write(dev: &mut SpiDevice, tx_buf: &[u8]) -> Result {
    Spi::write_then_read(dev, tx_buf, &mut [0u8; 0])
}

#[inline]
fn read(dev: &mut SpiDevice, rx_buf: &mut [u8]) -> Result {
    Spi::write_then_read(dev, &[0u8; 0], rx_buf)
}
```



Merging our work

A screenshot of a GitHub comment interface. The background is a blurred image of a squirrel in a tree. The comment is from user 'ksquirrel' and is 6 days old. The comment text is 'Review of 0861303f78de :'. Below this, there is a list of four commit checks, each with a green checkmark icon, the word 'Commit', a commit hash in a blue box, and the text ': Looks fine!'. The commit hashes are 7e65b95, cdefa5d, f650715, and 0861303.

ksquirrel commented 6 days ago

Review of 0861303f78de :

- ✓ Commit 7e65b95 : Looks fine!
- ✓ Commit cdefa5d : Looks fine!
- ✓ Commit f650715 : Looks fine!
- ✓ Commit 0861303 : Looks fine!

- [illegible]

[illegible]

Any questions?

- esteban.blanc@lse.epita.fr
- martin.schmidt@lse.epita.fr
- cohenarthur.dev@gmail.com

<https://github.com/ks0n/linux>
<https://github.com/ks0n/mfrc522-linux>

